

WEEK-12

SESSION-23

Task1: To Execute Large Pipelines, Launch Multiple Jobs, Measure Stability, Monitor Resources, Detect Failures, Analyze Performance.

```
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task2$ cd ~/OSSP/skill
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill$ mkdir -p week12/task1
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill$ cd week12/task1
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ nano pipeline_jobs.c
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ gcc pipeline_jobs.c -o
T1
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ ./T1
=====
          SKILL WEEK 12 - TASK 1
        LARGE PIPELINE TEST
=====

Launching 5 pipeline jobs...

Job 1 started. PID = 3082
Job 2 started. PID = 3083
Job 3 started. PID = 3086
Job 4 started. PID = 3088
Job 5 started. PID = 3089
99996
99997
99998
99999
100000
Job 1 completed successfully.
99996
99997
99998
99999
100000
99996
99997
99998
99999
100000
Job 3 completed successfully.
Job 2 completed successfully.
99996
99997
99998
```

```
Job 4 completed successfully.
sh: 1: invalid_command: not found
Job 5 failed.

=====
          PERFORMANCE REPORT
=====
Total Jobs      : 5
Completed Jobs  : 4
Failed Jobs     : 1
Total Time      : 0.3818 seconds
Stability Status : FAILURE DETECTED
=====
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ top
top - 10:43:27 up 3:08, 1 user, load average: 0.06, 0.03, 0.00
Tasks: 31 total, 1 running, 27 sleeping, 2 stopped, 1 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.
MiB Mem : 5778.6 total, 5013.2 free, 555.4 used, 353.9 buff/cache
MiB Swap: 2048.0 total, 2048.0 free, 0.0 used, 5223.2 avail Mem

  PID USER      PR  NI    VIRT    RES    SHR  S  %CPU  %MEM   TIME+
3124 madhuli+  20   0    8172    6120   4012  R   0.3   0.1   0:00.02
  1 root      20   0   24252   15368  11568  S   0.0   0.3   0:01.11
  2 root      20   0    3180    2208    2072  S   0.0   0.0   0:00.01
  6 root      20   0    3180    2048    1940  S   0.0   0.0   0:00.00
 46 root      19  -1   50400   16916  15604  S   0.0   0.3   0:00.46
 82 systemd+  20   0   22416   14488   11996  S   0.0   0.2   0:00.10
 90 root      20   0   35352   12296    9216  S   0.0   0.2   0:06.13
157 root      20   0    2888    2012    1884  S   0.0   0.0   0:00.04
159 root      20   0    4472    3140    2892  S   0.0   0.1   0:00.05
161 message+  20   0    8824    5320    4640  S   0.0   0.1   0:00.07
176 root      20   0   44100   29884   14748  S   0.0   0.5   0:00.19
181 root      20   0   18436    9400    8156  S   0.0   0.2   0:00.13
204 root      20   0    5492    2808    2620  S   0.0   0.0   0:00.00
212 syslog    20   0  220548   5904    4600  S   0.0   0.1   0:00.10
226 _chrony    20   0   20424    6424    5644  S   0.0   0.1   0:00.12
231 root      20   0  123464   32720   18112  S   0.0   0.6   0:00.19
250 _chrony    20   0   12096   2444    1796  S   0.0   0.0   0:00.00
339 root      20   0    3196    1116     980  S   0.0   0.0   0:00.00
340 root      20   0    3196    1252    1108  S   0.0   0.0   0:01.25
342 madhuli+  20   0    6268    5652   3928  S   0.0   0.1   0:00.40
```

```

madhuLiKa@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ ps -eo pid,ppid,stat,%c
pu,%mem,cmd --sort=-%cpu | head
  PID    PPID  STAT  %CPU  %MEM  CMD
  3030      1  Ssl   0.0   0.2  /usr/libexec/wsl-pro-service
    90      1  Ss    0.0   0.2  /usr/lib/systemd/systemd-udevd
   340     339  S     0.0   0.0  /init
    1       0  Ss    0.0   0.2  /sbin/init
   46      1  S<s   0.0   0.2  /usr/lib/systemd/systemd-journald
  342     340  Ss    0.0   0.0  -bash
  387      1  Ss    0.0   0.2  /usr/lib/systemd/systemd --user
  231      1  Ssl   0.0   0.5  /usr/bin/python3 /usr/share/unattended-upgrad
es/unattended-upgrade-shutdown --wait-for-signal
  176      1  Ss    0.0   0.5  /usr/bin/python3 /usr/bin/networkd-dispatcher
--run-startup-triggers
madhuLiKa@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ seq 1 1000000 | grep '[
0-9]' | sort -n | tail -n 10
999991
999992
999993
999994
999995
999996
999997
999998
999999
1000000
madhuLiKa@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ time sh -c "seq 1 10000
00 | grep '[0-9]' | sort -n | tail -n 10"
999991
999992
999993
999994
999995
999996
999997
999998
999999
1000000

real    0m0.228s
user    0m0.171s

```

```

Job 4 completed successfully.
99996
99997
99998
99999
100000
=====
          SKILL WEEK 12 - TASK 1
          LARGE PIPELINE TEST
=====

Launching 5 pipeline jobs...

Job 1 started. PID = 3242
Job 1 completed successfully.
sh: 1: invalid_command: not found
=====
          SKILL WEEK 12 - TASK 1
          LARGE PIPELINE TEST
=====

Launching 5 pipeline jobs...

Job 5 started. PID = 3249
Job 5 failed.
=====
          SKILL WEEK 12 - TASK 1
          LARGE PIPELINE TEST
=====

Launching 5 pipeline jobs...

=====
          PERFORMANCE REPORT
=====
Total Jobs      : 5
Completed Jobs  : 4
Failed Jobs     : 1
Total Time      : 0.3083 seconds
Stability Status : FAILURE DETECTED

```

```

GNU nano 8.7.1 pipeline_jobs.c *
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/time.h>

#define NUM_JOBS 5

double get_time()
{
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return tv.tv_sec + tv.tv_usec / 1000000.0;
}

int main()
{
    pid_t pids[NUM_JOBS];
    double start_time, end_time;
    int status;
    int completed = 0;
    int failed = 0;

    printf("=====\n");
    printf("          SKILL WEEK 12 - TASK 1\n");
    printf("          LARGE PIPELINE TEST\n");
    printf("=====\n\n");

    printf("Launching %d pipeline jobs...\n\n", NUM_JOBS);

    start_time = get_time();

    for (int i = 0; i < NUM_JOBS; i++)
    {
        pids[i] = fork();

        if (pids[i] < 0)

```

```

        exit(0);
    }

    for (int i = 0; i < NUM_JOBS; i++)
    {
        if (pids[i] > 0)
        {
            waitpid(pids[i], &status, 0);

            if (WIFEXITED(status) && WEXITSTATUS(status) == 0)
            {
                completed++;
            }
            else
            {
                failed++;
            }
        }
    }

    end_time = get_time();

    printf("\n=====\n");
    printf("          PERFORMANCE REPORT\n");
    printf("=====\n");

    printf("Total Jobs      : %d\n", NUM_JOBS);
    printf("Completed Jobs   : %d\n", completed);
    printf("Failed Jobs      : %d\n", failed);
    printf("Total Time       : %.4f seconds\n",
        end_time - start_time);

    if (failed == 0)
    {
        printf("Stability Status : STABLE\n");
    }
}

```

```
GNU nano 8.7.1 monitor.sh *
#!/bin/bash

echo "=====
echo "      RESOURCE MONITORING REPORT"
echo "=====

echo ""
echo "CPU and Memory Usage:"
ps -eo pid,ppid,stat,%cpu,%mem,cmd --sort=-%cpu | head -10

echo ""
echo "System Load:"
uptime

echo ""
echo "Memory Information:"
free -h

echo ""
echo "Process Count:"
ps -e --no-headers | wc -l

echo ""
echo "=====
```

Observation: The large pipeline execution and multiple job management system was successfully implemented and tested. Multiple jobs were launched simultaneously using separate processes, and the execution time of the pipelines was measured. The system monitored process and resource information such as CPU and memory usage using Linux monitoring commands. Successful jobs were completed normally, while an intentionally failed pipeline was detected and reported. The performance reports were generated to analyze execution time, stability, and resource usage.

Task2: To Document Architecture, Create Flow Diagrams, Explain Execution Paths, Describe Syscalls, Record Design Decisions, Review Documentation.

```
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task1$ cd ~/OSSP/skill/week12
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12$ mkdir -p task2
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12$ cd task2
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task2$ nano architecture_documentation.md
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task2$ cat architecture_documentation.md
# SKILL WEEK 12 - TASK 2
```

```
## Documentation of Pipeline Job Execution System
```

```
---
```

1. Objective

To document the architecture, create flow diagrams, explain execution paths, describe system calls, record design decisions, and review the documentation of the Linux pipeline job execution system.

```
---
```

2. System Overview

The system is designed to execute multiple Linux pipeline jobs using process creation and process management techniques.

The system launches multiple child processes using `fork()`, executes pipeline commands, waits for child processes using `waitpid()`, detects successful and failed jobs, and measures the execution time.

Linux monitoring commands are also used to observe CPU, memory, process, and system resource usage.

```
---
```

3. System Architecture

The system consists of the following major components:

1. Main Controller
2. Process Creation Module

```
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task2$ nano syscall_table.txt
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task2$ cat architecture_documentation.md
# SKILL WEEK 12 - TASK 2
```

```
## Documentation of Pipeline Job Execution System
```

```
---
```

1. Objective

To document the architecture, create flow diagrams, explain execution paths, describe system calls, record design decisions, and review the documentation of the Linux pipeline job execution system.

```
---
```

2. System Overview

The system is designed to execute multiple Linux pipeline jobs using process creation and process management techniques.

The system launches multiple child processes using `fork()`, executes pipeline commands, waits for child processes using `waitpid()`, detects successful and failed jobs, and measures the execution time.

Linux monitoring commands are also used to observe CPU, memory, process, and system resource usage.

```
---
```

3. System Architecture

The system consists of the following major components:

1. Main Controller
2. Process Creation Module
3. Pipeline Execution Module
4. Job Management Module
5. Failure Detection Module
6. Performance Measurement Module

```
GNU nano 8.7.1 architecture_documentation.md *
# SKILL WEEK 12 - TASK 2

## Documentation of Pipeline Job Execution System

---

## 1. Objective

To document the architecture, create flow diagrams, explain execution paths, describe system calls, record design decisions, and review the documentation of the Linux pipeline job execution system.

---

## 2. System Overview

The system is designed to execute multiple Linux pipeline jobs using process creation and process management techniques.

The system launches multiple child processes using fork(), executes pipeline commands, waits for child processes using waitpid(), detects successful and failed jobs, and measures the execution time.


Linux monitoring commands are also used to observe CPU, memory, process, and system resource usage.

---

## 3. System Architecture

The system consists of the following major components:

1. Main Controller
2. Process Creation Module
3. Pipeline Execution Module
4. Job Management Module
5. Failure Detection Module
```




13. Documentation Review Checklist

- [✓] System objective documented
- [✓] Architecture documented
- [✓] Execution flow documented
- [✓] Process execution path explained
- [✓] Important system calls described
- [✓] Resource monitoring explained
- [✓] Failure detection documented
- [✓] Performance measurement documented
- [✓] Design decisions recorded
- [✓] Advantages and limitations documented

14. Conclusion

The architecture and execution flow of the Linux pipeline job execution system were documented successfully. The process creation, pipeline execution, job monitoring, failure detection, performance measurement, and resource monitoring mechanisms were explained. Important Linux system calls and functions were identified, and the major design decisions were recorded. The documentation provides a clear description of how the system operates and how its components interact.





GNU nano 8.7.1 syscall_table.txt *	
=====	
SYSTEM CALL / FUNCTION TABLE	
=====	
Function	Purpose

fork()	Creates a new child process
waitpid()	Waits for a specific child process and obtains its termination status
gettimeofday()	Obtains current time for performance measurement
system()	Executes a Linux shell command
exit()	Terminates a process and returns status
top	Monitors CPU, memory and processes
ps	Displays process information
free	Displays memory usage
uptime	Displays system load information
time	Measures command execution time

Observation: The architecture documentation for the Linux pipeline job execution system was successfully prepared. The system architecture, execution flow, process execution path, system calls, resource monitoring, failure detection, performance measurement, and design decisions were documented. A flow diagram was created to represent the complete execution path, and a system-call table was prepared to describe the important Linux functions used by the system. The documentation was reviewed using the provided checklist to ensure that all major components and execution steps were covered.

SESSION-24

Task1: To Organize Source Code, Improve Build Scripts, Add README, Verify Repository Structure, Tag Releases, Validate Deliverables.

```
madhulika@LAPTOP-D2M8NKBH:~/OSSP/skill/week12/task2$ cd ~/OSSP
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ git status
On branch skill
Your branch is up to date with 'origin/skill'.

Untracked files:
  (use "git add <file>..." to include in what will be committed)
    Week1/
    Week3/Task2/filecopy
    interactive_loop
    interactive_loop.c
    practical/
    report.md
    skill/
    tokenizer.c

nothing added to commit but untracked files present (use "git add" to track)
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ git branch
  Practical
  Practical-backup
  main
* skill
  skill-sessions
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ git remote -v
origin  https://github.com/Madhulika024/OSSP.git (fetch)
origin  https://github.com/Madhulika024/OSSP.git (push)
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ find . -maxdepth 3 -type f | sort
./.git/COMMIT_EDITMSG
./.git/FETCH_HEAD
./.git/HEAD
./.git/ORIG_HEAD
./.git/config
./.git/description
./.git/hooks/applypatch-msg.sample
./.git/hooks/commit-msg.sample
./.git/hooks/fsmonitor-watchman.sample
./.git/hooks/post-update.sample
./.git/hooks/pre-applypatch.sample
./.git/hooks/pre-commit.sample
```

```
- Git
- Linux system calls and process-management APIs

## Build

C programs can generally be compiled using:

gcc source.c -o program

The generated executable can then be run using:

./program

## Testing

Programs are tested using different input conditions and Linux
process-management commands. Outputs are verified against the expected
results.

## Skills Covered

The repository covers process creation, process synchronization,
signals, pipes, FIFOs, process groups, job control, error handling,
pipeline execution, resource monitoring, testing, and documentation.

## Validation

Source files, executable programs, reports, and documentation are
reviewed before final submission.

## Version Control

Git is used to maintain the source code history and organize project
deliverables. # OSSP
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ nano Makefile
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ make
OSSP Build System
Use gcc source.c -o program to build individual programs.
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ make clean
find . -type f -executable -delete
Executable files removed.
```

```
GNU nano 8.7.1                                README.md *
# OSSP - Operating System and Shell Programming

## Overview

This repository contains the Operating System and Shell Programming
skills and practical exercises completed using Linux and C programming.

## Objectives

- Understand Linux process management
- Work with system calls
- Implement process creation and synchronization
- Understand signals and process groups
- Implement pipes and pipelines
- Study job control
- Handle errors and runtime failures
- Monitor processes and system resources
- Test and analyze Linux programs

## Repository Structure

OSSP/
├── skill/
│   ├── week1/
│   ├── week2/
│   ├── week3/
│   ├── week4/
│   ├── week5/
│   ├── week6/
│   ├── week7/
│   ├── week8/
│   ├── week9/
│   ├── week10/
│   ├── week11/
│   └── week12/
└──
```

```
GNU nano 8.7.1                                README.md *
├── practical/
│   ├── week1/
│   ├── week2/
│   ├── week3/
│   └── ...
└──

## Technologies Used

- Linux / Ubuntu
- C
- Bash Shell
- GCC
- Git
- Linux system calls and process-management APIs

## Build

C programs can generally be compiled using:

gcc source.c -o program

The generated executable can then be run using:

./program

## Testing

Programs are tested using different input conditions and Linux
process-management commands. Outputs are verified against the expected
results.

## Skills Covered

The repository covers process creation, process synchronization,
signals, pipes, FIFOs, process groups, job control, error handling,
pipeline execution, resource monitoring, testing, and documentation.

^G Help      ^O Write Out  ^F Where Is  ^K Cut       ^T Execute
^X Exit      ^R Read File  ^\ Replace   ^U Paste     ^J Justify
```

```
GNU nano 8.7.1                                Makefile *
```

```
CC = gcc
CFLAGS = -Wall -Wextra -g

.PHONY: all clean

all:
    @echo "OSSP Build System"
    @echo "Use gcc source.c -o program to build individual programs."

clean:
    find . -type f -executable -delete
    @echo "Executable files removed."
```

Observation: The project repository structure, source code organization, build configuration, documentation files, and supporting directories were reviewed. The required project components such as source files, header files, scripts, tests, documentation, and supporting resources were identified and verified. The repository was examined to ensure that the project deliverables were organized and available for final validation.

Task2: To Create Demo Scenarios, Validate Features, Rehearse Presentation, Demonstrate Pipelines, Demonstrate Signals, Final Project Review.

```
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ cd ~/OSSP
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ find . -type f -executable | sort
madhulika@LAPTOP-D2M8NKBH:~/OSSP$ cd /tmp
madhulika@LAPTOP-D2M8NKBH:/tmp$ cat > pipeline_demo.c <<'EOF'
#include <stdio.h>
#include <unistd.h>

int main() {
    printf("Pipeline Demo Started\n");

    printf("Process 1: Data generation\n");
    sleep(1);

    printf("Process 2: Data processing\n");
    sleep(1);

    printf("Process 3: Data output\n");
    sleep(1);

    printf("Pipeline execution completed successfully\n");

    return 0;
}
EOF
madhulika@LAPTOP-D2M8NKBH:/tmp$ gcc pipeline_demo.c -o pipeline_demo
madhulika@LAPTOP-D2M8NKBH:/tmp$ ./pipeline_demo
Pipeline Demo Started
Process 1: Data generation
Process 2: Data processing
Process 3: Data output
Pipeline execution completed successfully
madhulika@LAPTOP-D2M8NKBH:/tmp$ cat > signal_demo.c <<'EOF'
#include <stdio.h>
#include <signal.h>
#include <unistd.h>

void handle_signal(int sig) {
    if (sig == SIGUSR1)
        printf("SIGUSR1 received successfully\n");
    else if (sig == SIGTERM)
        printf("SIGTERM received successfully\n");
```

```

        printf("SIGTERM received successfully\n");
    }
}

int main() {
    signal(SIGUSR1, handle_signal);
    signal(SIGTERM, handle_signal);

    printf("Signal Demo Started\n");
    printf("Process ID: %d\n", getpid());
    printf("Waiting for signals...\n");

    sleep(2);

    raise(SIGUSR1);
    sleep(1);

    raise(SIGTERM);

    printf("Signal handling completed successfully\n");

    return 0;
}
EOF
madhulika@LAPTOP-D2M8NKBH:/tmp$ gcc signal_demo.c -o signal_demo
madhulika@LAPTOP-D2M8NKBH:/tmp$ ./signal_demo
Signal Demo Started
Process ID: 3751
Waiting for signals...
SIGUSR1 received successfully
SIGTERM received successfully
Signal handling completed successfully
madhulika@LAPTOP-D2M8NKBH:/tmp$ cat > process_demo.c <<'EOF'
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    pid_t pid = fork();

    if (pid < 0) {
        perror("fork");

```

```

    }

    if (pid == 0) {
        printf("Child Process Started\n");
        printf("Child PID: %d\n", getpid());
        sleep(2);
        printf("Child Process Completed\n");
    } else {
        printf("Parent Process Started\n");
        printf("Parent PID: %d\n", getpid());
        printf("Created Child PID: %d\n", pid);

        wait(NULL);

        printf("Parent: Child process completed\n");
        printf("Process management demonstration completed\n");
    }

    return 0;
}
EOF
madhulika@LAPTOP-D2M8NKBH:/tmp$ cat > process_demo.c <<'EOF'
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    pid_t pid = fork();

    if (pid < 0) {
        perror("fork");
        return 1;
    }

    if (pid == 0) {
        printf("Child Process Started\n");
        printf("Child PID: %d\n", getpid());
        sleep(2);

```

```

madhulika@LAPTOP-D2M8NKBH:/tmp$ echo "=====
echo "          OSSP PROJECT DEMO"
echo "=====

echo ""
echo "1. PIPELINE DEMONSTRATION"
./pipeline_demo

echo ""
echo "2. SIGNAL DEMONSTRATION"
./signal_demo

echo ""
echo "3. PROCESS MANAGEMENT DEMONSTRATION"
./process_demo

echo ""
echo "=====
echo "          FINAL PROJECT REVIEW"
echo "=====
echo "Pipeline execution: Demonstrated"
echo "Signal handling: Demonstrated"
echo "Process management: Demonstrated"
echo "Parent-child synchronization: Demonstrated"
echo "Demo review: Completed"
=====
          OSSP PROJECT DEMO
=====

1. PIPELINE DEMONSTRATION
Pipeline Demo Started
Process 1: Data generation
Process 2: Data processing
Process 3: Data output
Pipeline execution completed successfully

2. SIGNAL DEMONSTRATION
Signal Demo Started
Process ID: 3863
Waiting for signals...
SIGUSR1 received successfully

```

```

SIGUSR1 received successfully
SIGTERM received successfully
Signal handling completed successfully

3. PROCESS MANAGEMENT DEMONSTRATION
Parent Process Started
Parent PID: 3879
Created Child PID: 3880
Child Process Started
Child PID: 3880
Child Process Completed
Parent: Child process completed
Process management demonstration completed

=====
          FINAL PROJECT REVIEW
=====
Pipeline execution: Demonstrated
Signal handling: Demonstrated
Process management: Demonstrated
Parent-child synchronization: Demonstrated
Demo review: Completed

```

Observation: The final project demonstration was successfully conducted by executing separate demonstrations for pipeline processing, signal handling, and process management. The pipeline demonstration showed sequential processing, the signal demonstration showed signal generation and handling, and the process demonstration showed parent-child process creation and synchronization. The complete demonstration was reviewed successfully.

Result:

The pipeline, signal handling, process management, and parent-child synchronization demonstrations were successfully executed and verified.